



Arm CCA Open Source Enablement Status

21-June-2024, Kevin Zhao, Linaro

Agenda

CCA Software Stack Introduction

CCA attestation

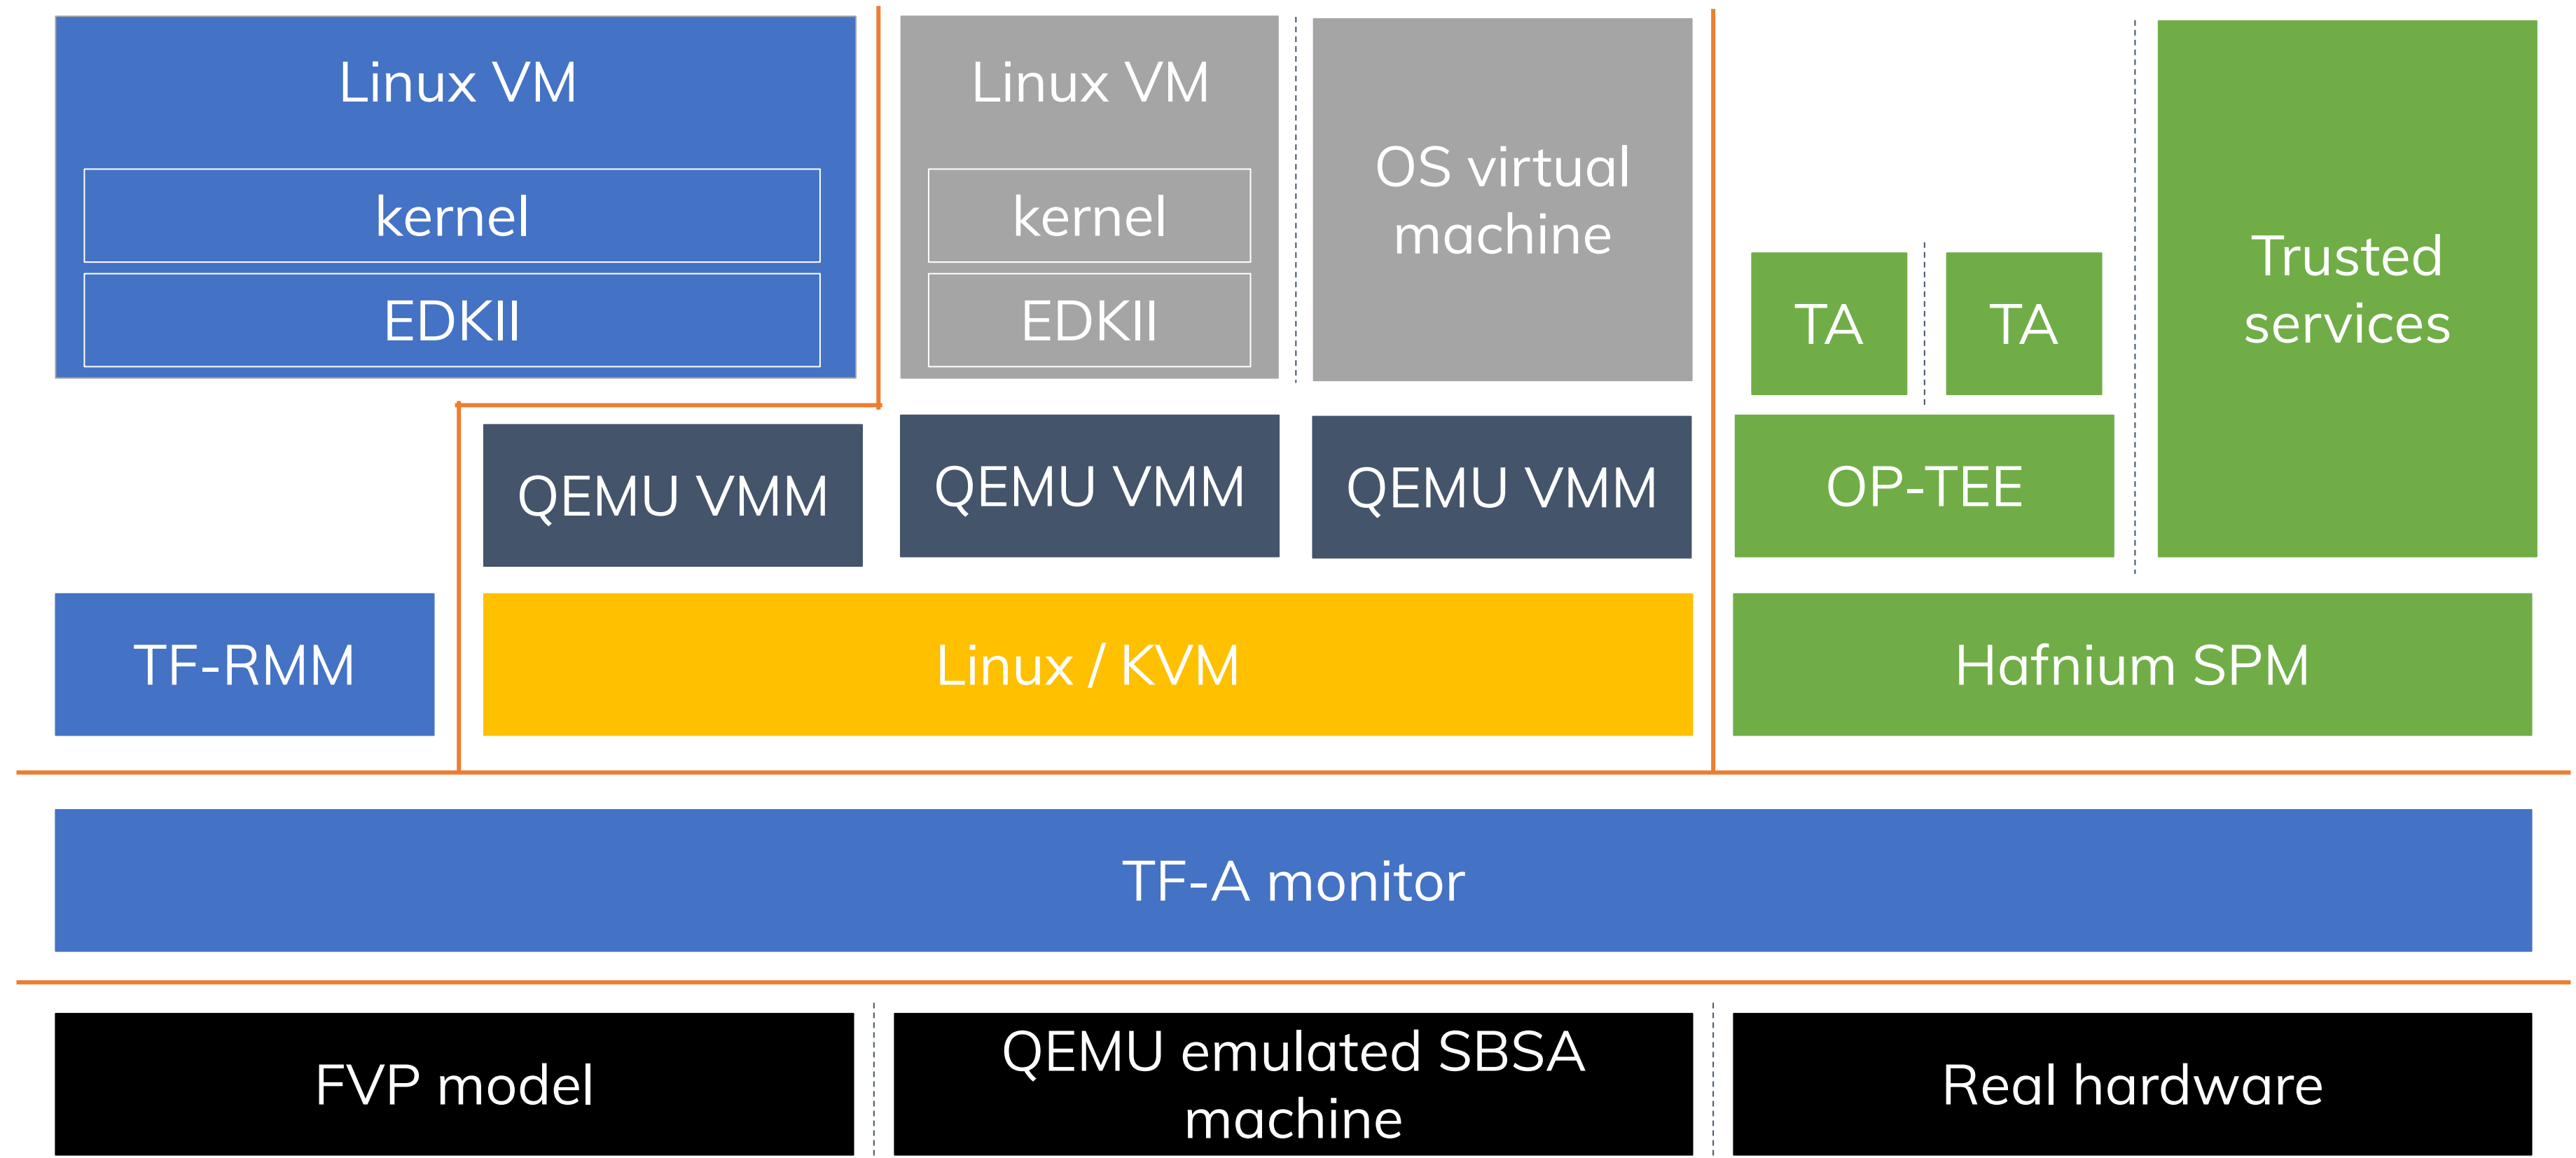
CCA 1.1 Specs

CCA Software Stack Introduction



Arm Solutions at Lightspeed

Low level reference software stack



CCA Low level software reference stack

Arm developed a CCA stack that runs on their FVP model.

With the release of QEMU 8.1, Linaro ported that stack to QEMU:

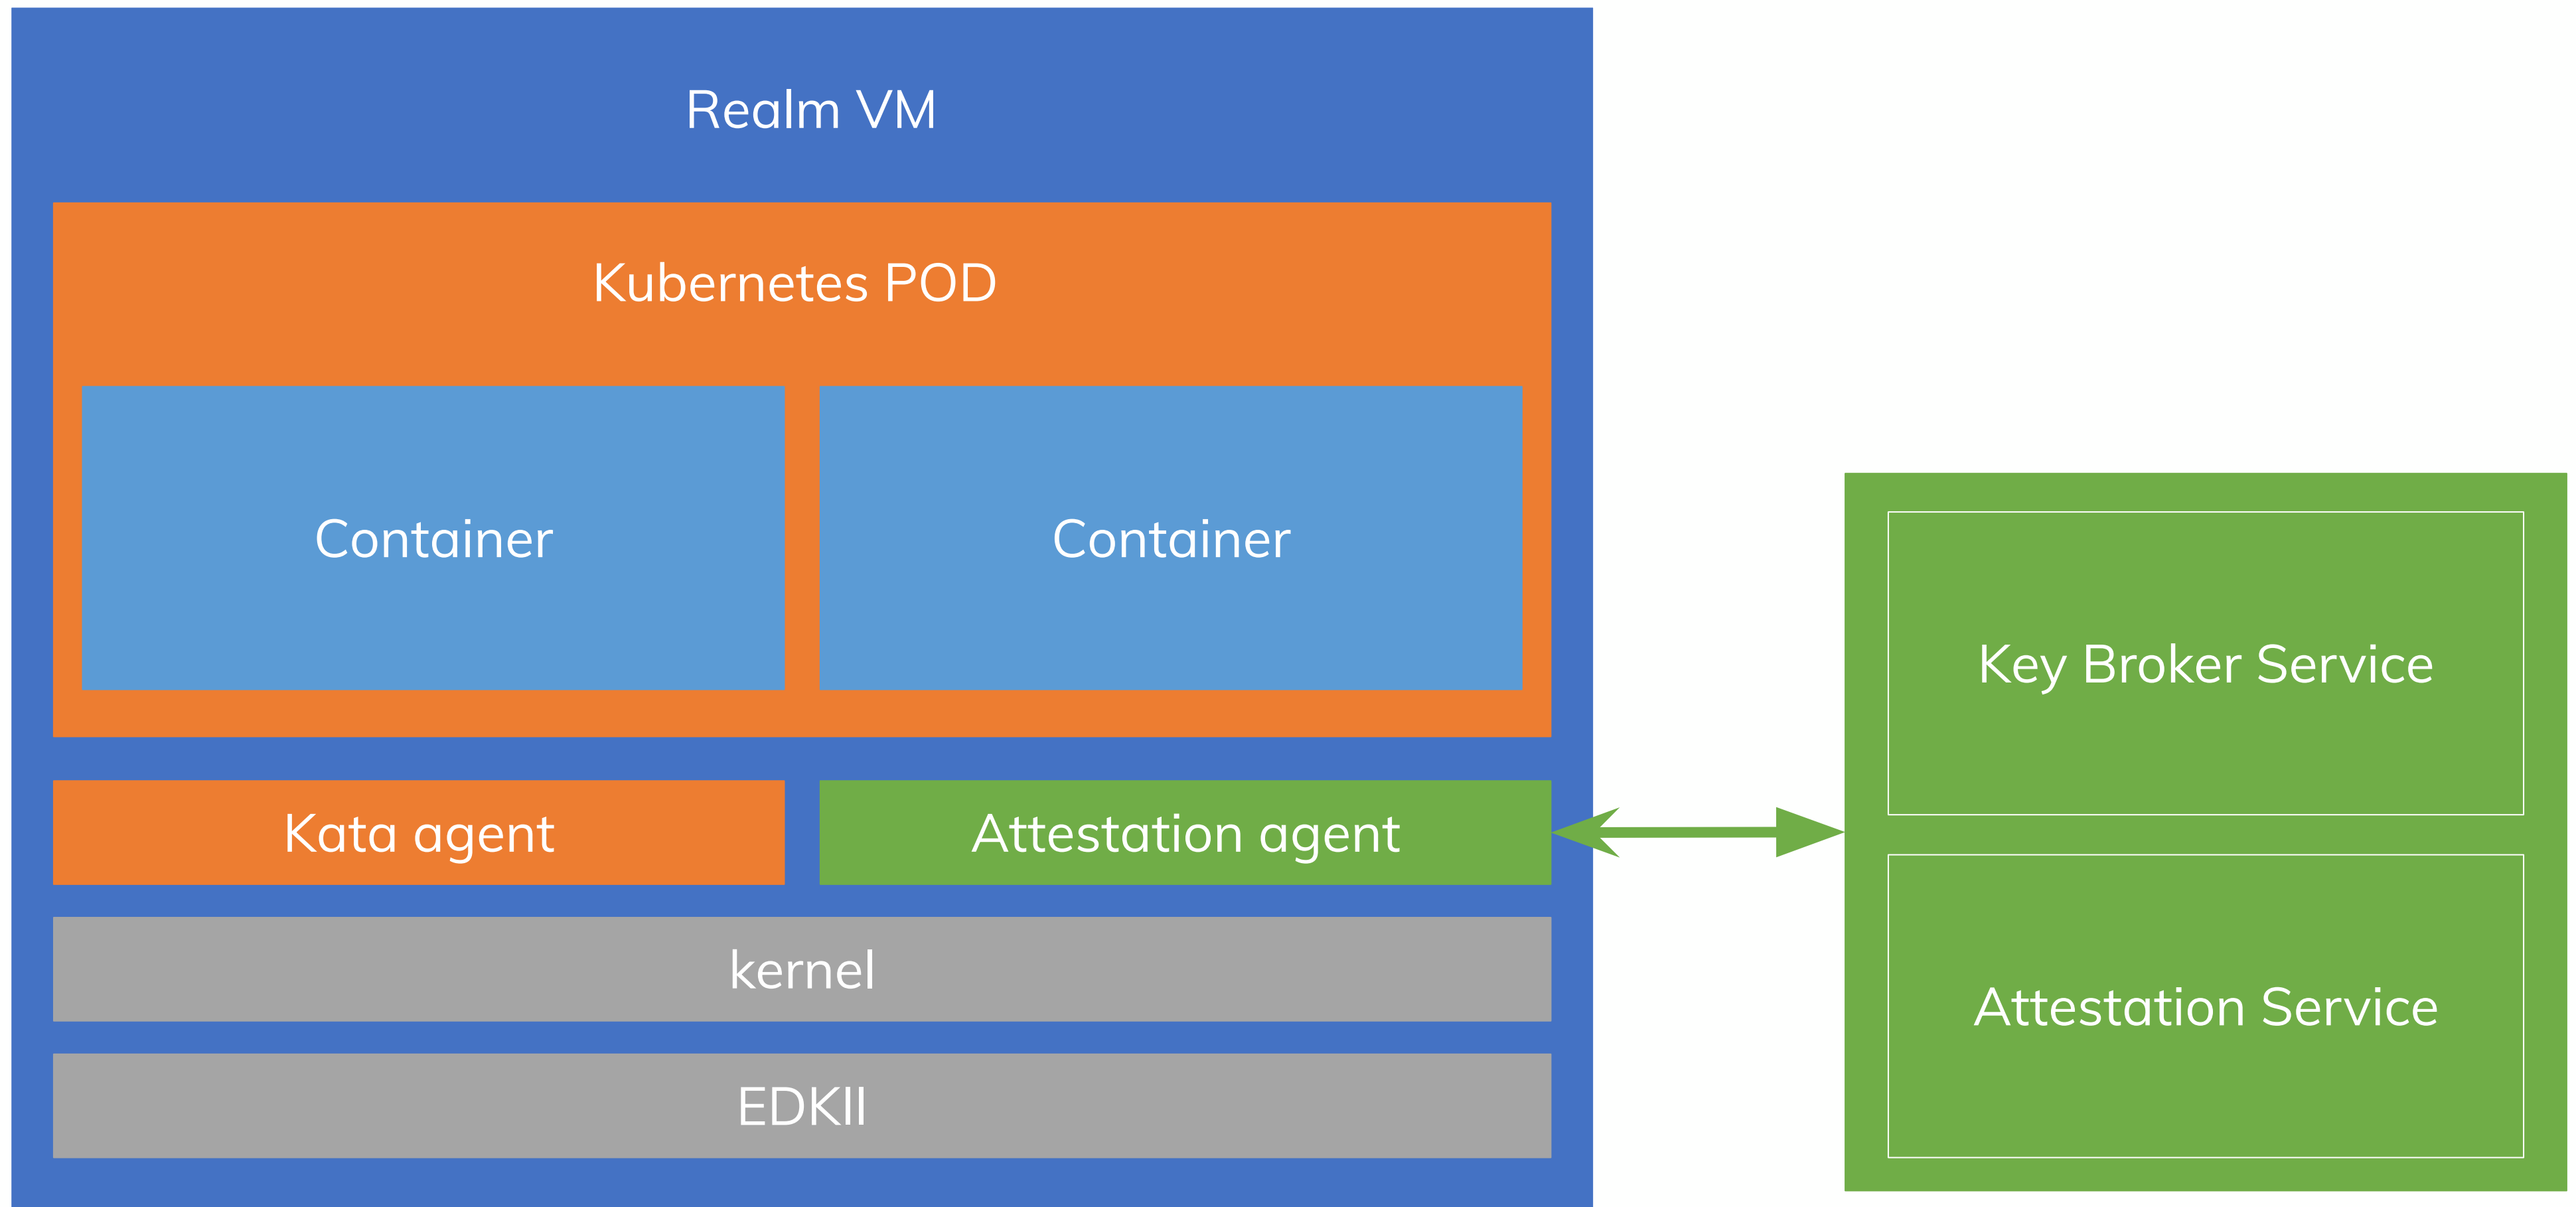
- Patches for TF-A, TF-RMM and EDK2 are available on [CodeLinaro](#) (cca/v2 branch).
- Patches for the [Linux kernel](#) and [kvmtool](#) are hosted by Arm (cca/v2 branch).
- The solution is currently for the QEMU virt machine type with buildroot.
 - QEMU as a system emulator with RME support and as a VMM launching Realms.

Work to support QEMU SBSA reference machine type is ongoing.

Support for RME in Linaro's Trusted Reference Stack (TRS) is ongoing. Plans for a CI.

[Documentation](#) is available to compile and run the stack, from base system to Realm.

High level reference software stack



CCA high level software reference stack

Kata container support: [code repo](#).

- Current features:
 - Only supports Kata v2.
 - Only supports QEMU back end.
 - Only supports direct kernel boot with Kata. The UEFI boot disk image has been validated.
 - Only supports ACPI=off in QEMU.

Confidential Containers (CoCo):

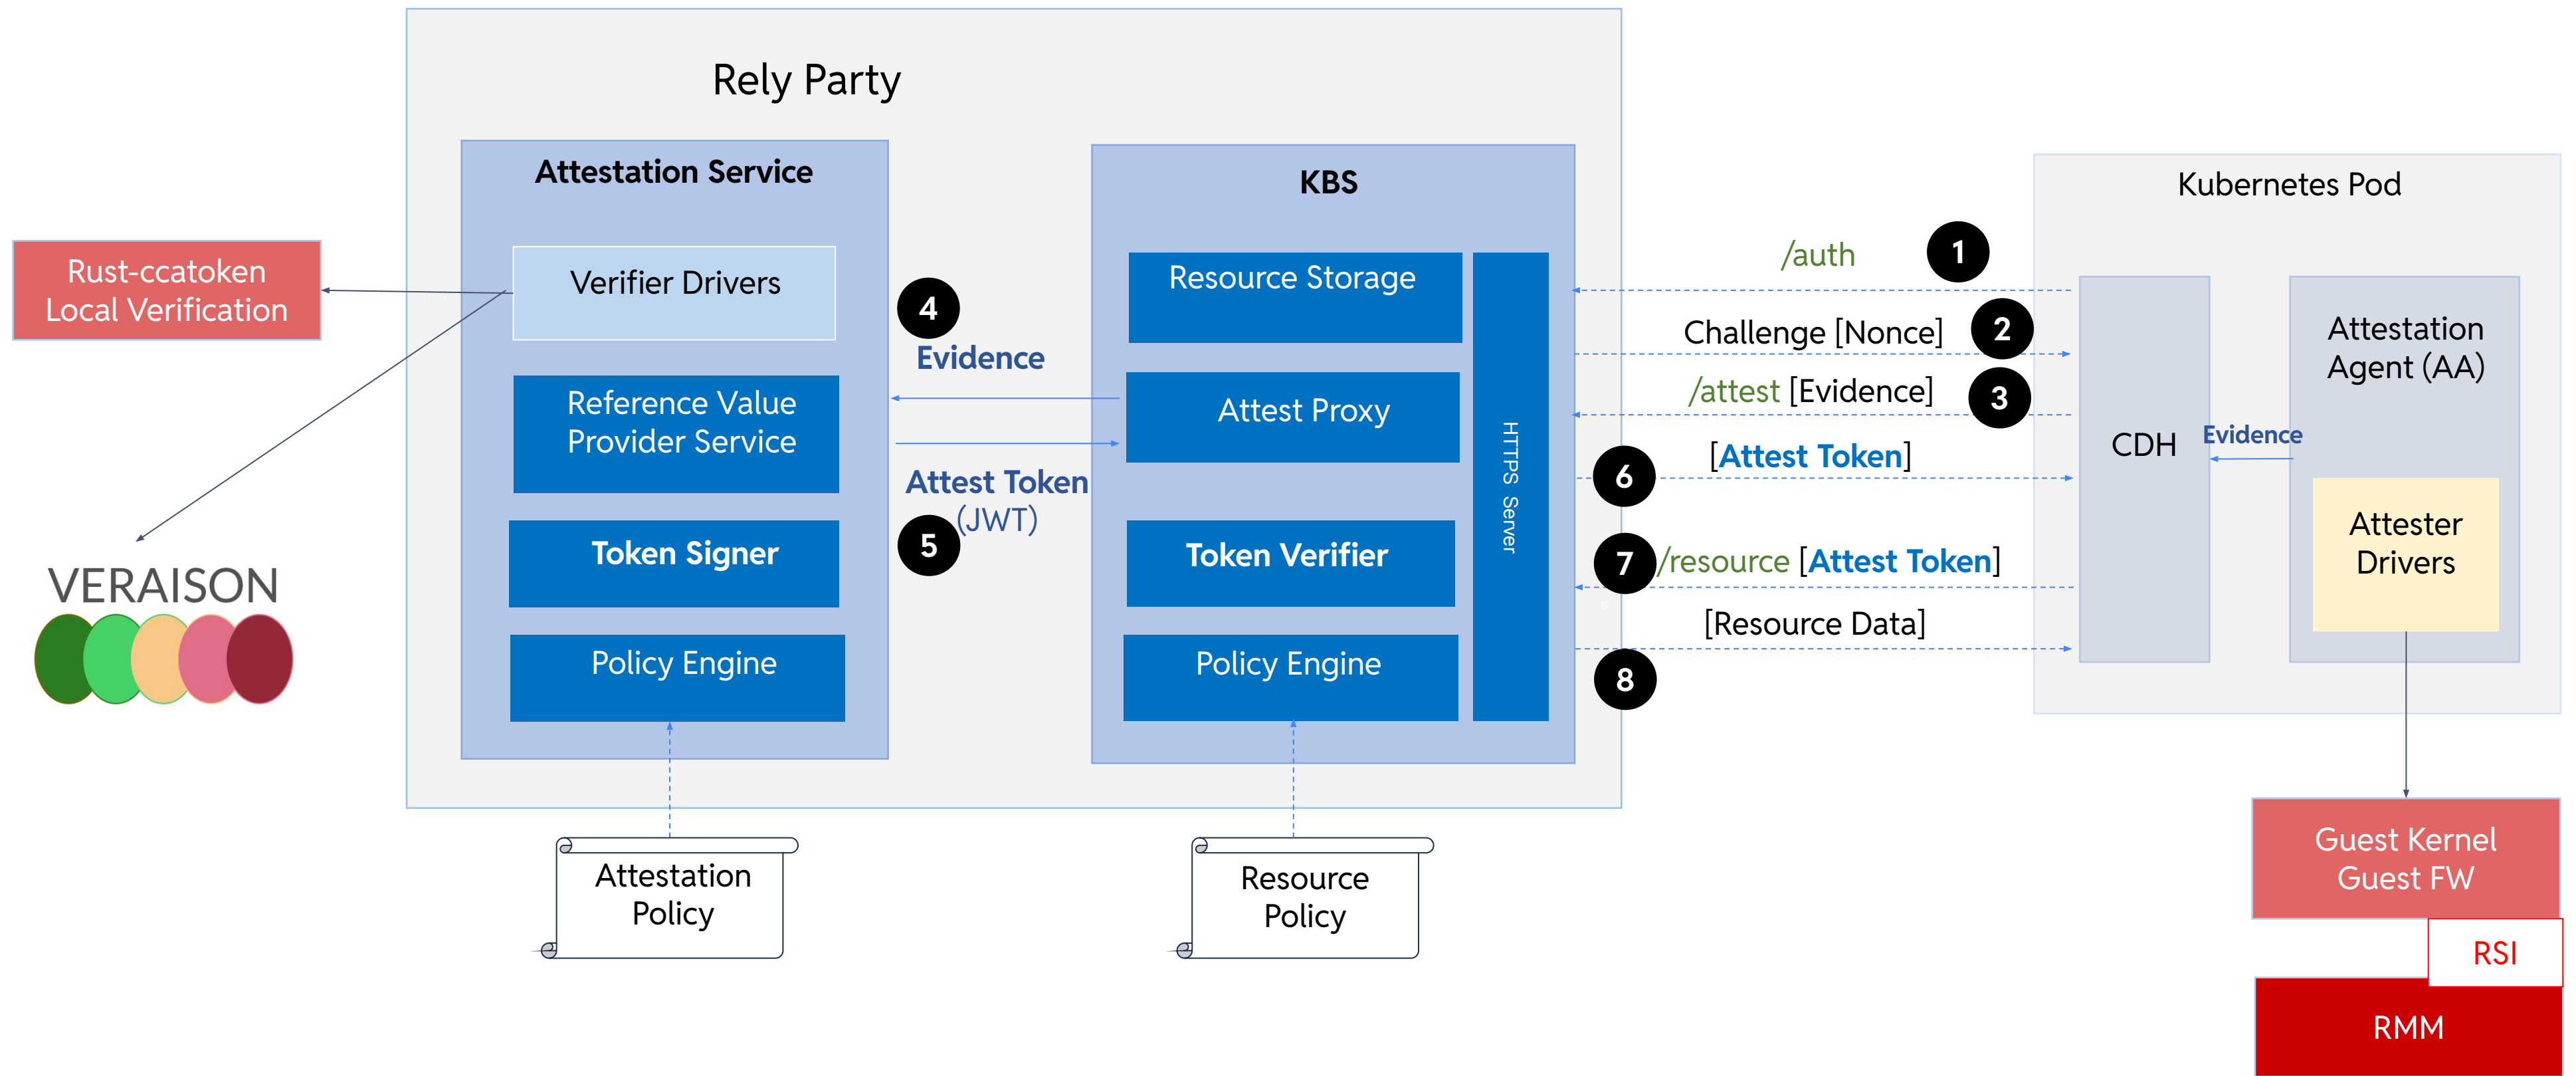
- Framework adoption
 - Kubernetes Confidential Computing operator
 - Container image service (service offload, encryption verification).
- Trustee support: [code repo](#).

CCA Attestation



Arm Solutions at Lightspeed

CoCo and Veraison - remote attestation



Attestation tools

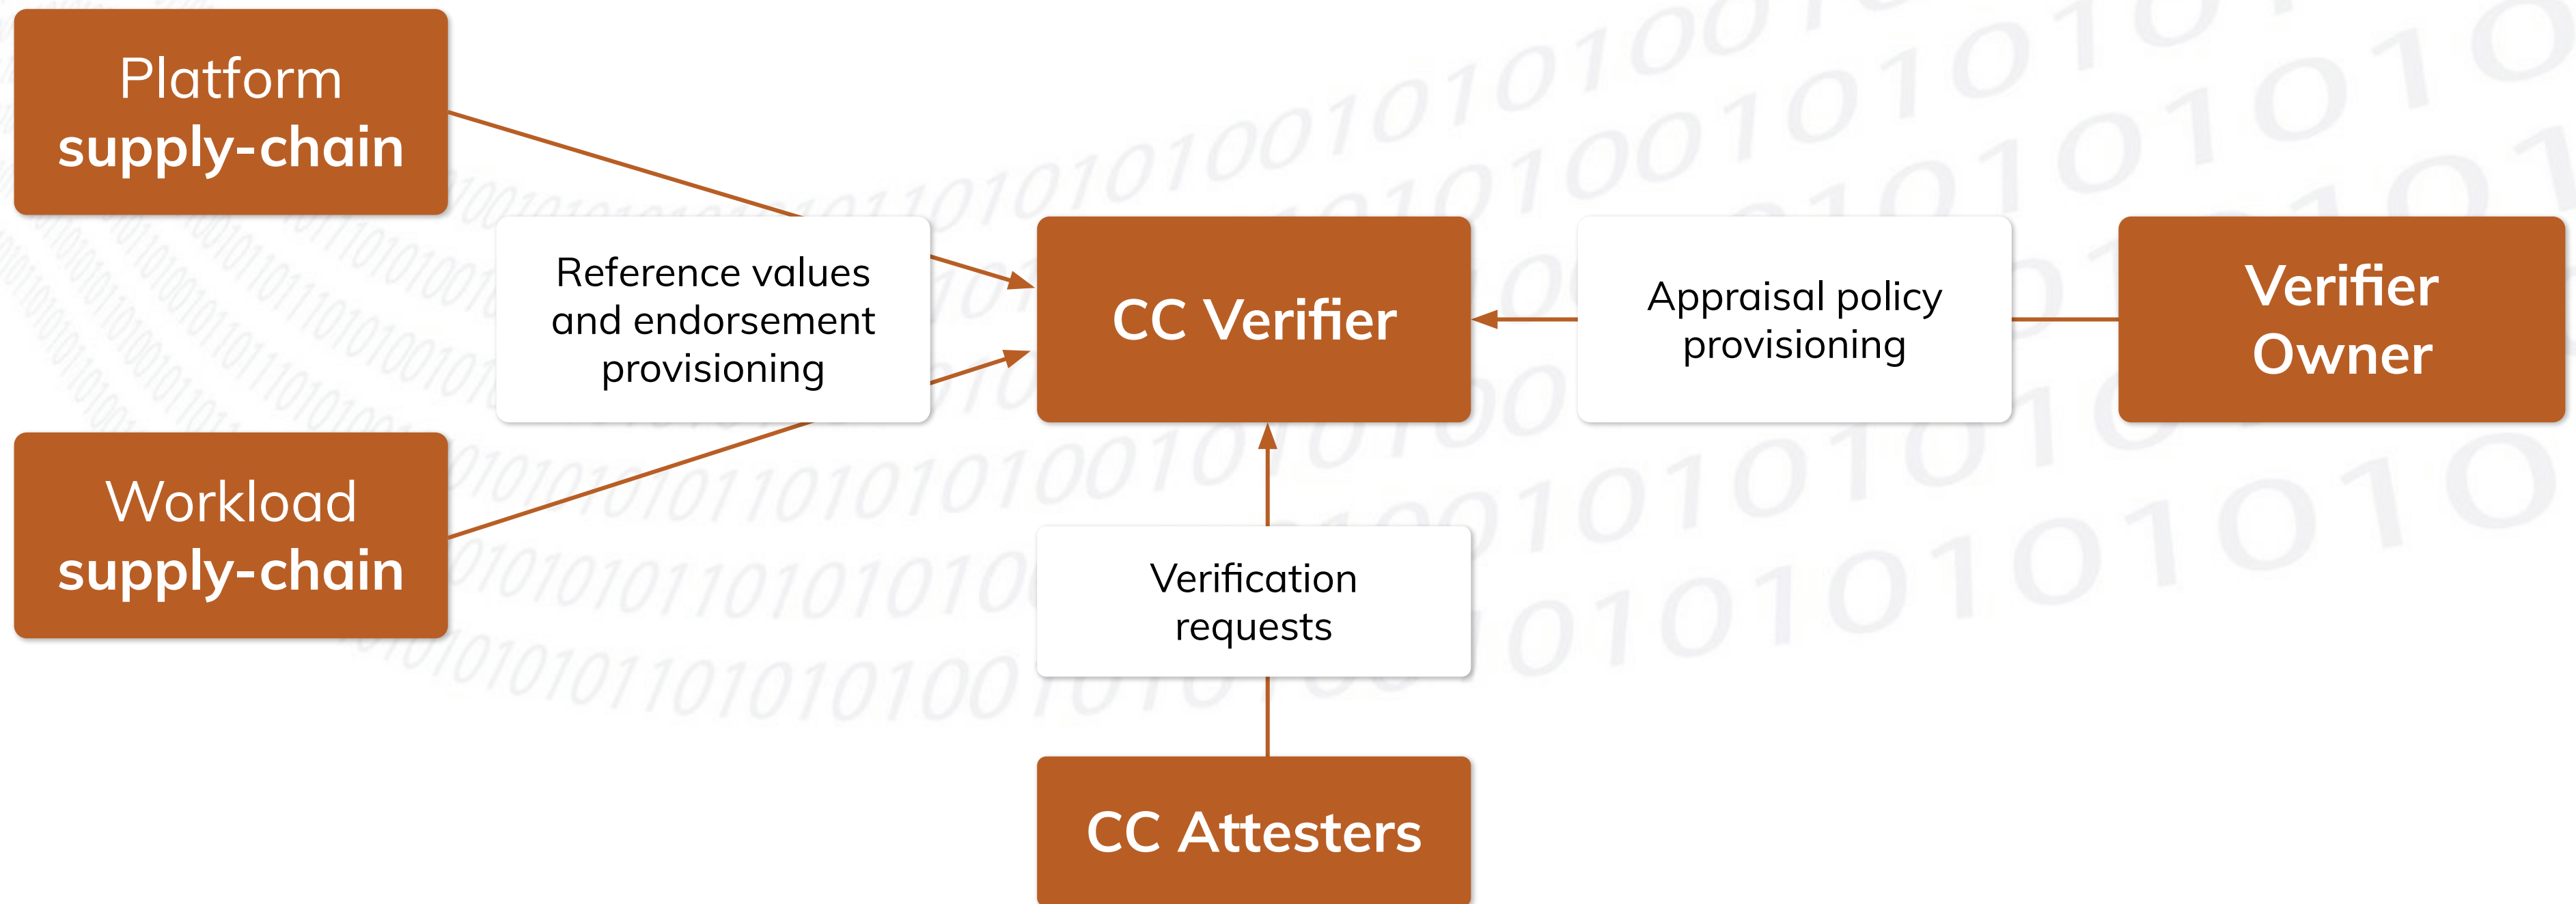
ccatoken crate

- Provides command line tools and APIs to decode and verify CCA attestation tokens.
- Published at <https://crates.io/crates/ccatoken>.
- Sources at <https://github.com/veraison/rust-ccatoken>.

realm-token crate

- Tool that calculates the Realm initial and extended measurements, needed for CCA attestation.
 - Sources at <https://git.codelinaro.org/linaro/dcap/realm-token>.

Remote attestation verification



Future steps

QEMU support for memory encryption.

QEMU support for SMMU. This is a requirement for device assignment.

Cloud Hypervisor support.

Lightweight firmware support for Arm CCA.

End-to-end demo for CoCo on Arm CCA with Qemu backend.

Integration of Veraison and CoCo Attestation Service (AS) to provide a holistic end to end reference solution for confidential containers on Arm platforms.

CCA 1.1 Specs



Arm Solutions at Lightspeed

CCA 1.1 features – needed for initial deployments

Further strengthen the security guarantees provided to end users (Realm owners)

- Memory Encryption Contexts (MEC)
 - Physical memory contents of each Realm protected using a unique key or tweak
- Multiple signers
 - Require firmware image to be endorsed by multiple authorities, for example vendor plus a trusted auditor


Enable migration of workloads from non-secure VM to Realm, by providing feature parity

- Planes
 - Multiple privilege levels within a Realm, orthogonal to traditional kernel / user-space split
- Device Assignment (DA)
 - Enable trusted device functions to be admitted into a Realm's TCB, and granted DMA
- Host Debug of Realm
 - In a controlled environment, enable host to debug a realm (bypassing CCA security guarantees)

Allow platform owners additional flexibility, in deploying and updating firmware

- Live firmware activation
 - Update firmware image(s) while workloads continue to run, with minimal loss of availability
 - Replace platform firmware (for example, RMM) with an image supplied by the non-secure host

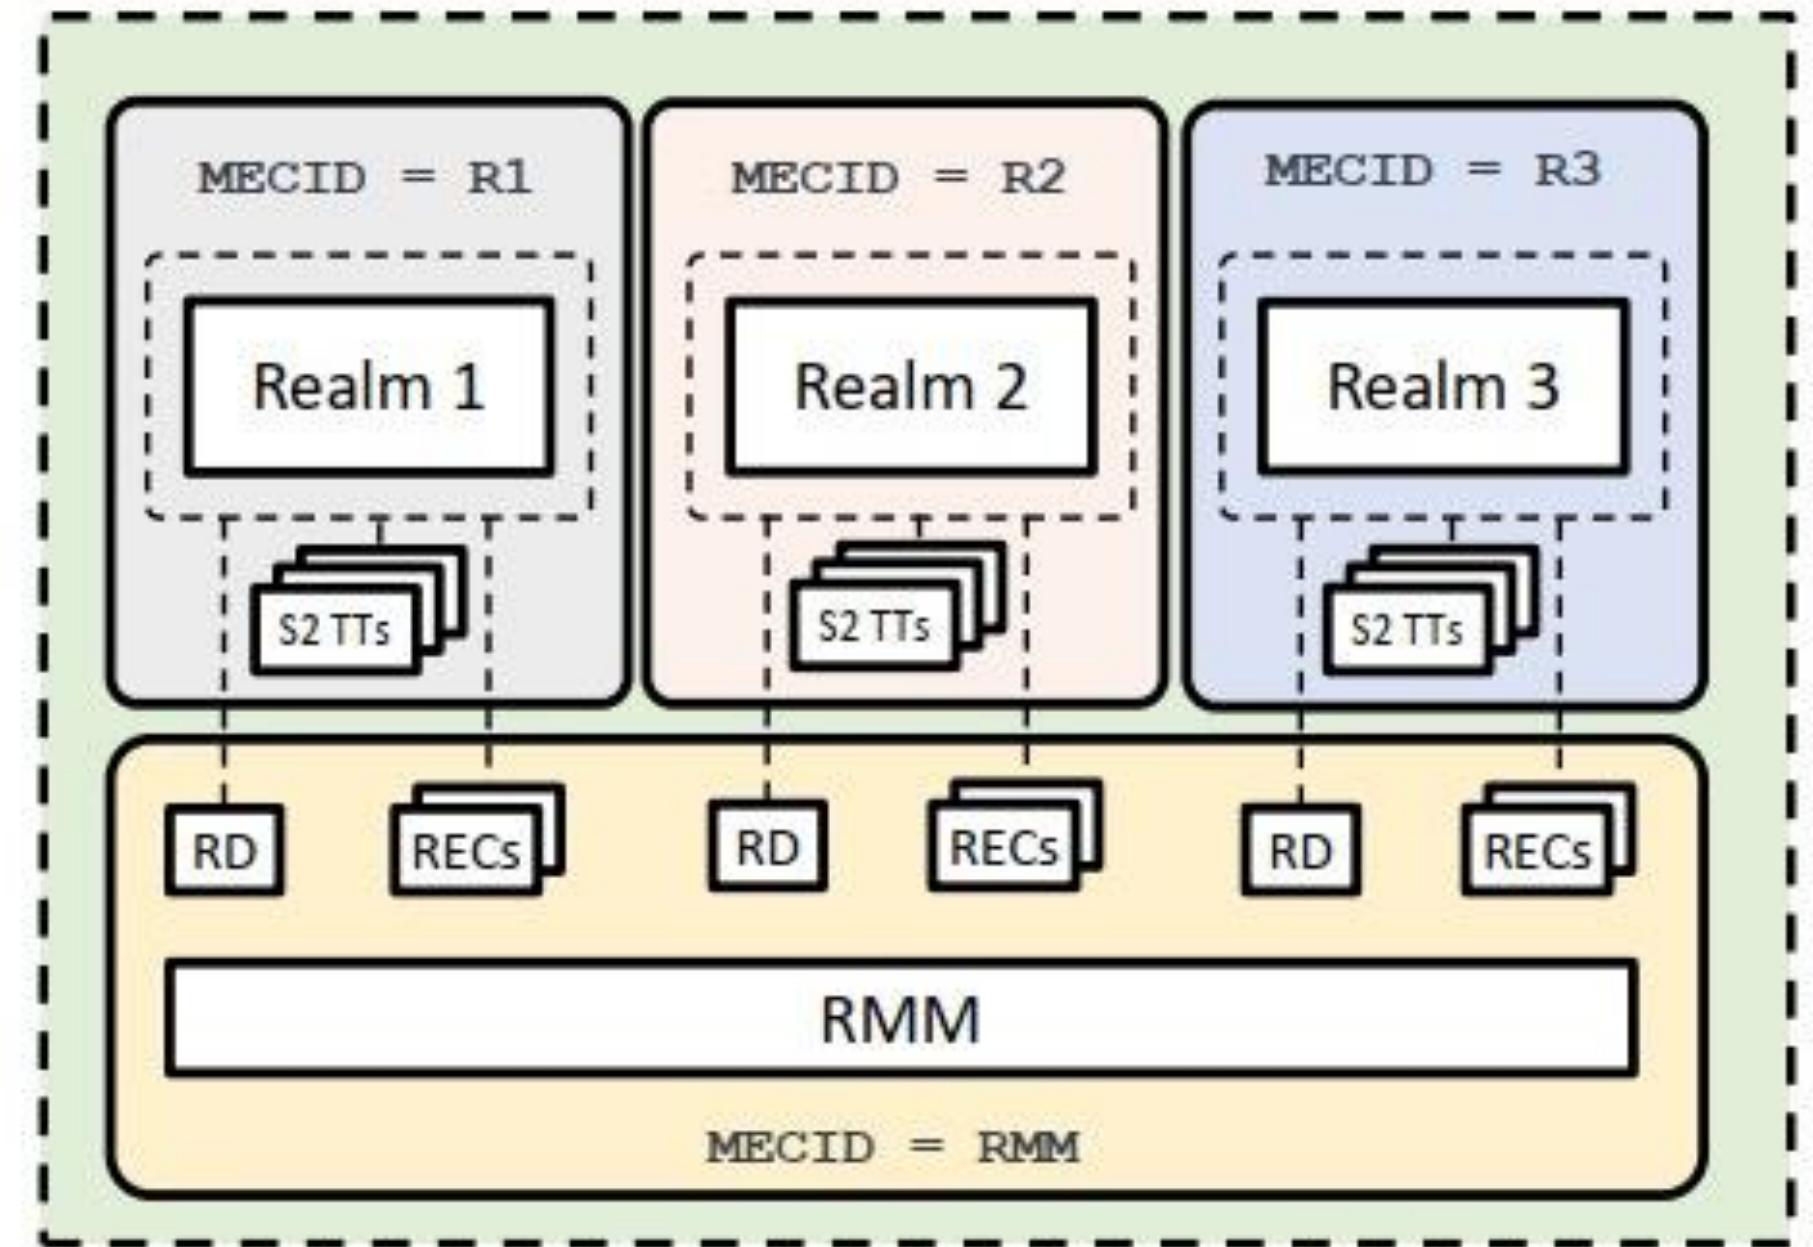
Arm CCA 1.0



- TF-RMM impact: Small
 - Enforce MEC ID uniqueness

Isolation Boundary
Crypto Boundary

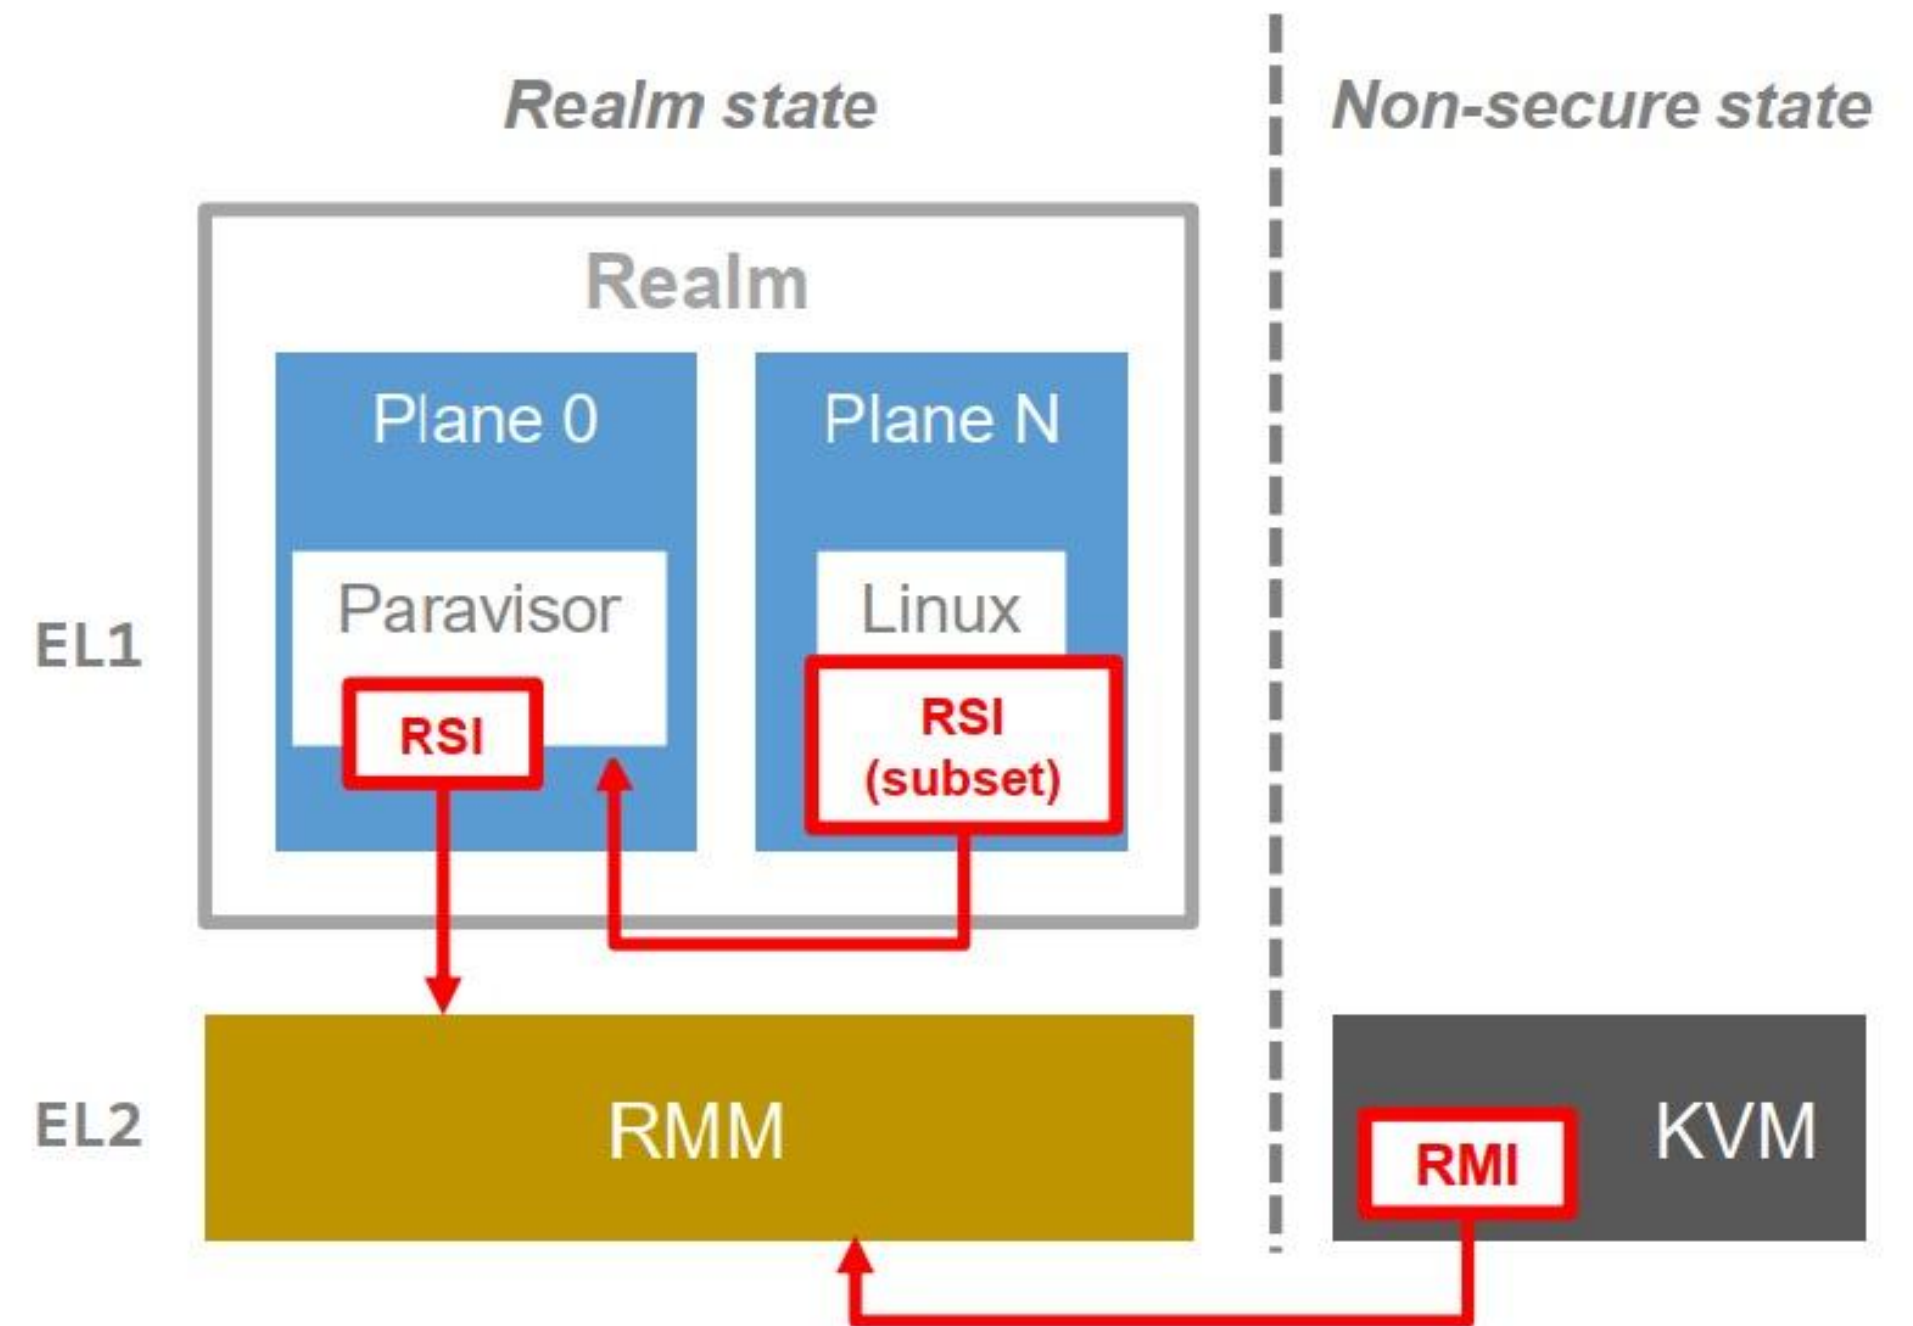
+ MEC



- Realm guest Linux impact: Zero
- Host Linux/KVM impact: X-Small
 - Allocate a MECID for each Realm

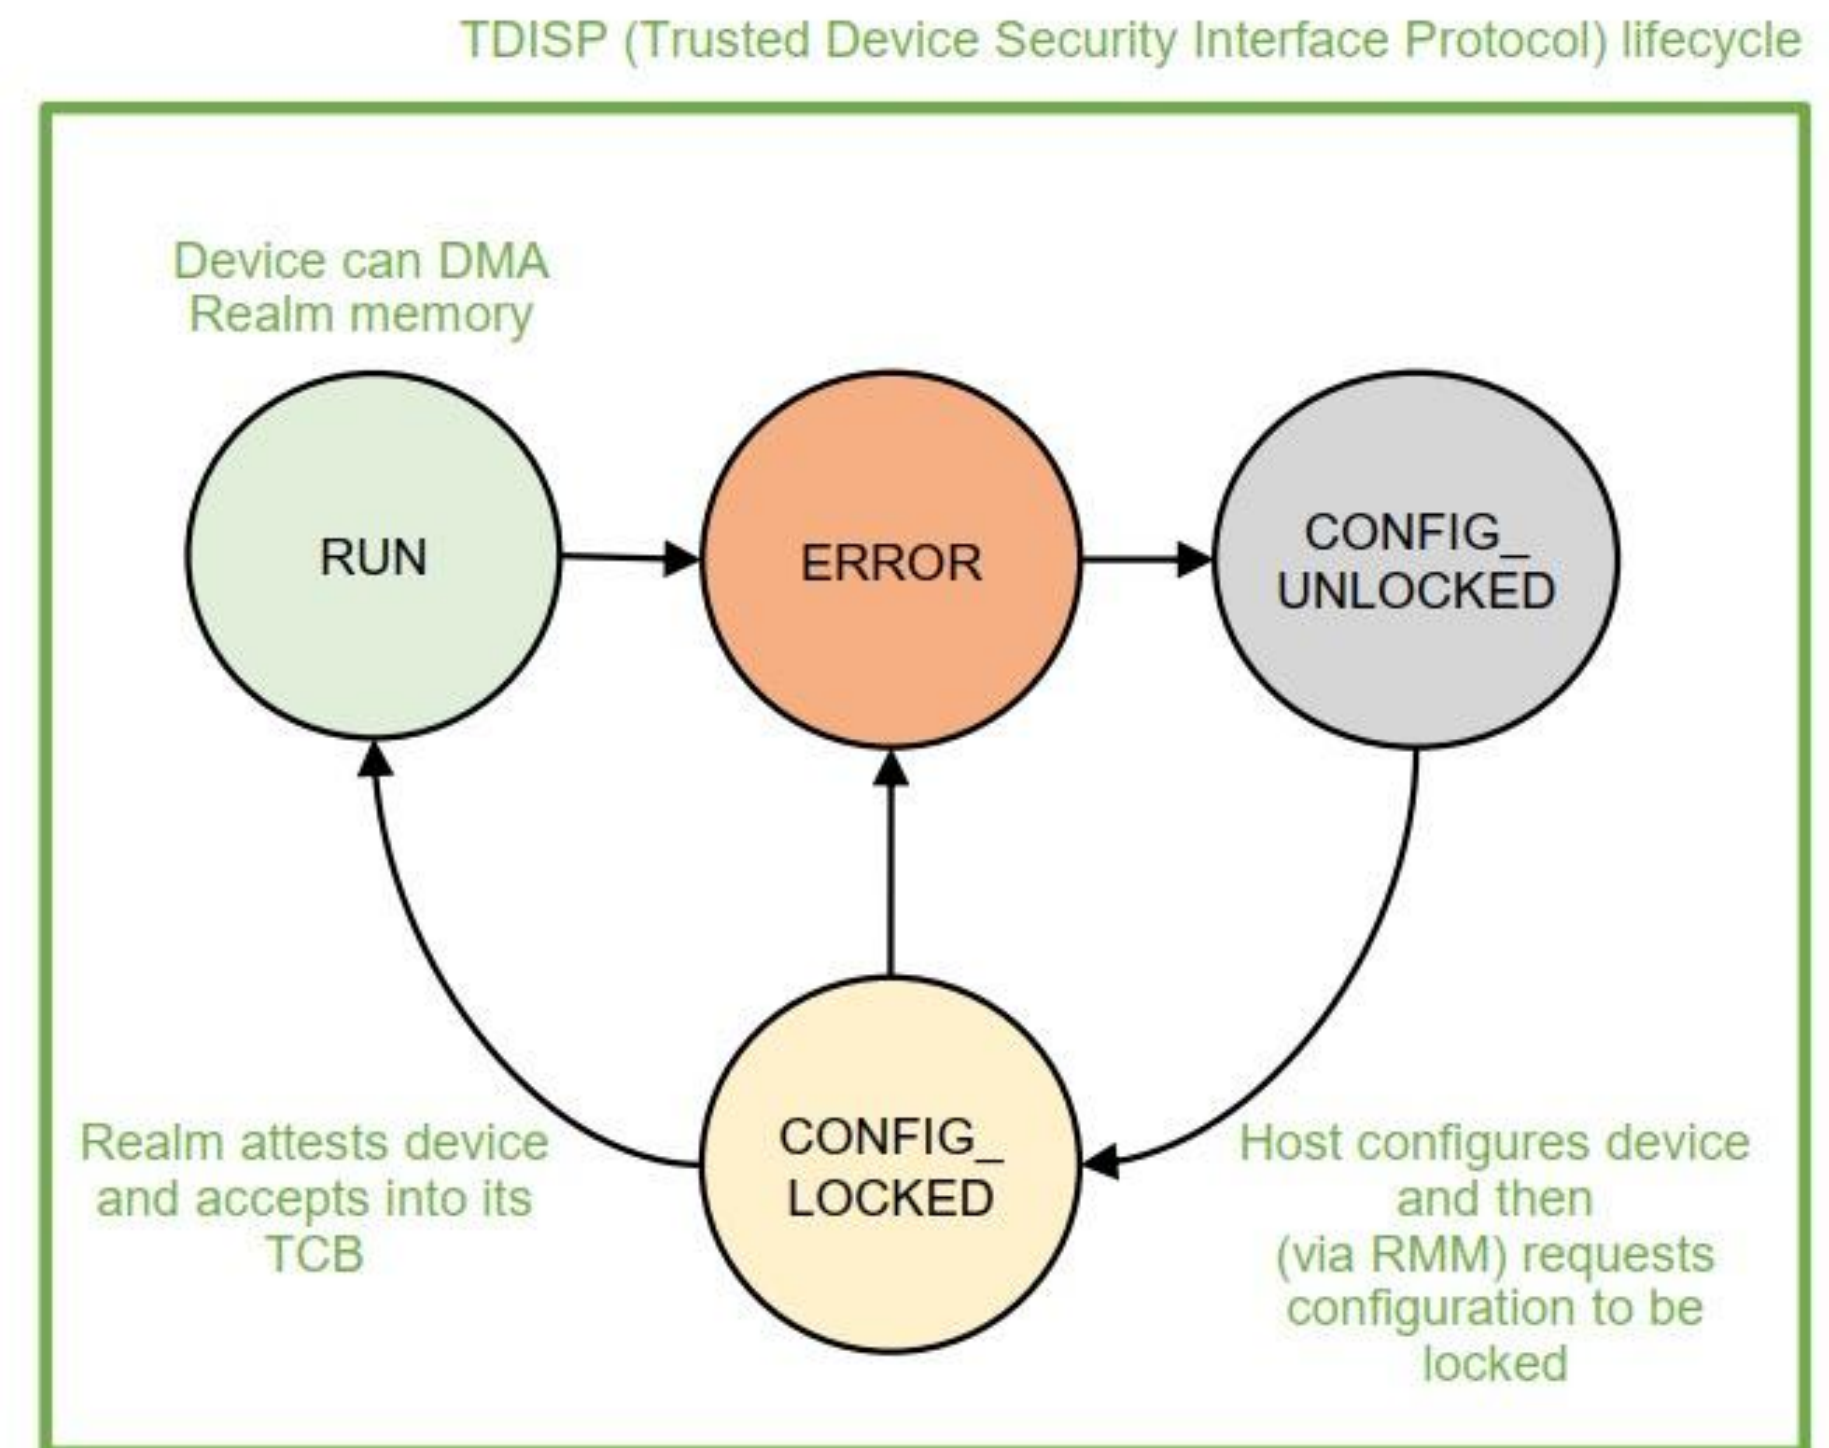
Planes

- In addition to the main guest OS and user workload, allow the contents of a Realm to include other software components (Planes)
 - For example, a security service like vTPM
- Provide isolation within a Realm, allowing privilege separation between the Planes
 - All Planes have Same IPA → PA mappings but IPA memory permissions may differ Allow the host hypervisor to continue treating
- the Realm as a single unit, for the purposes of resource allocation, scheduling and migration Within the
- Realm, privileged Plane 0 assigns resources to the other Planes



Device Assignment

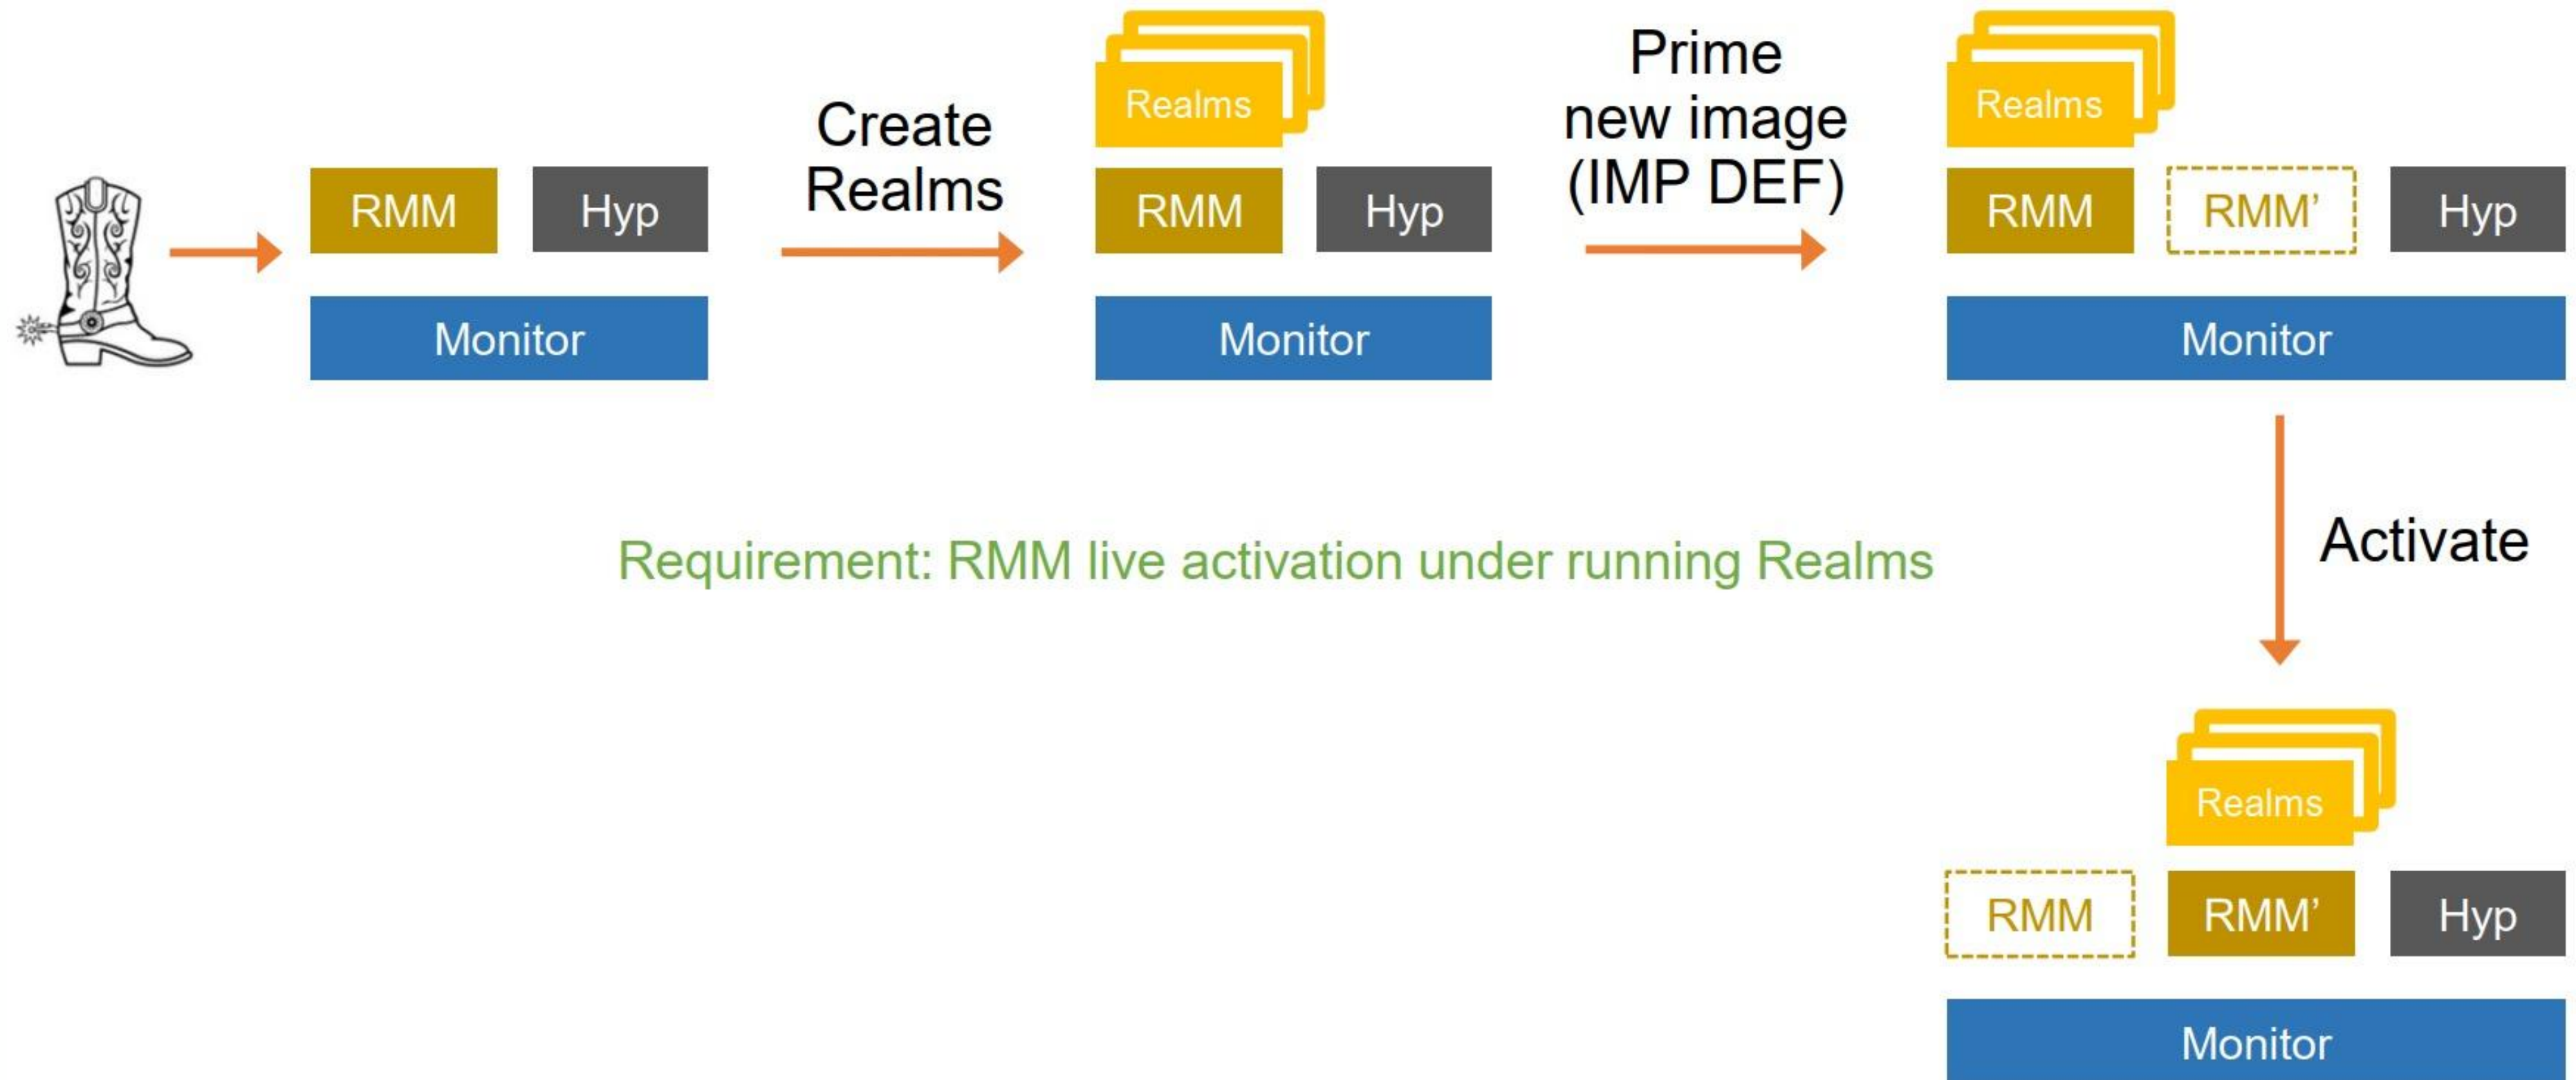
- Allow hypervisor to assign a PCIe TDISP device to a Realm
 - Also support coherently-attached devices, such as CXL instances*
 - Also support on-chip PCIe devices*
- Allow Realm to attest the identity and configuration of the device function
- Device lifecycle guarantees that
 - DMA is blocked until device has been approved by the Realm
 - Any changes in device configuration cause transition to an error state, which revokes DMA
 - Once removed from a Realm, device guarantees that it will scrub confidential state
- Management of device lifecycle must be standards-based
 - RMM must not require any device-specific knowledge
 - However, RMM will require knowledge of platform topology



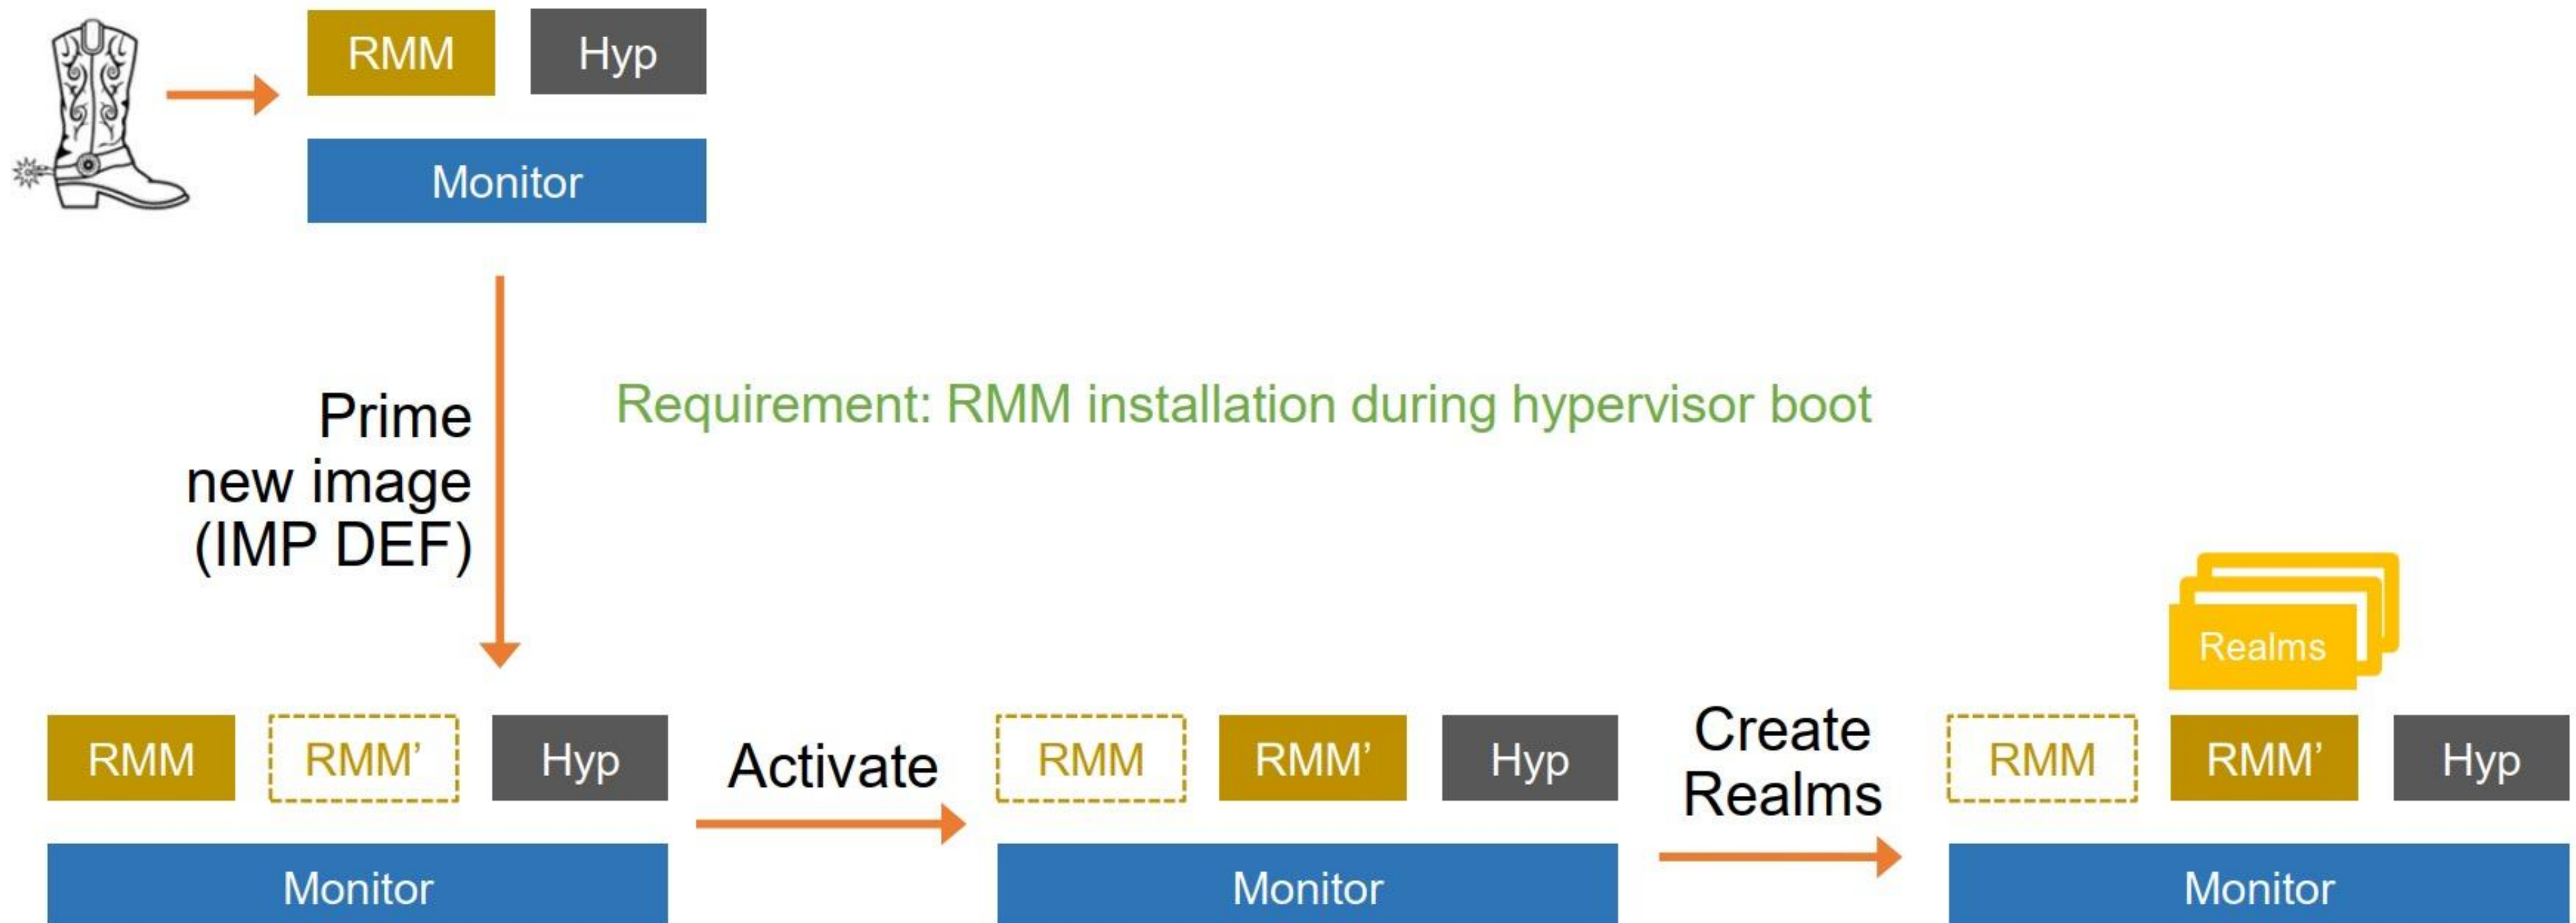
Device Assignment

- TF-A / TF-RMM impact: Large (foundation) -> X-Large (advanced)
 - Implement DA ABIs and integrate PCIe standard reference libs to enforce device lifecycle
 - SMMU S2 driver + SMMU S1 emulation + PCIe root port programming
- Realm guest Linux impact: Large (foundation) -> X-Large (advanced)
 - Generic PCIe / driver support for trusted devices
 - Irrespective of whether running on Confidential-Compute (CC) or non-CC VMs
 - Community Drive:
 - Driven by PCI Dev maintainer
 - Arch specific backends for TDISP will plug in to CCA interface hooks
- Host Linux/KVM impact: Large (foundation) -> X-Large (advanced)
 - DA enlightenment (use new RMIs to control device lifecycle)
 - UABI for VMM to describe DA devices – shared across architectures
 - SMMU stub (for interrupt management) for RMM SMMU Driver

Live firmware activation



Live firmware activation



Live firmware activation impacts

- TF-A / TF-RM
 - Create staging area for new firmware and transfer live state to new image
 - Live activating an arbitrary firmware version is hard – may need to restrict use-cases initially
 - For example, limit to specific code sections or require new version to be data compatible
 - Will focus on RMM and BL31 (EL3 firmware) live activate initially
 - Live activating the latter is especially hard (for example, may require CPU reset)
- Can increase use-cases over time (for example, by versioning data structures)
- Also need hooks to authenticate new firmware and to update firmware measurement log
 - Actual authentication is platform specific
- Realm guest Linux impact: None (hopefully)
- Host Linux/KVM impact: Medium
 - Use new ABIs to provide cycles to prime/activate new firmware
 - May need to quiesce activity and rendezvous CPUs during activation phase

Live firmware activation impacts

What to expect next

2024

- Continued upstreaming of CCA v1.0 Linux / KVM patches
- Monthly releases of RMM v1.1 spec
 - ALP with early DA / Planes support available now
 - Individual features will reach BET through the year, as they mature
- Collaborative development of CCA v1.1 SW (prototyping in progress)

2025

- Final RMM v1.1 spec (EAC)
- Upstreaming of non-DA-related CCA v1.1 features as they mature
 - Much quicker for TF projects than Linux / KVM
- Continued development of CCA v1.1 DA features and start upstreaming foundation support
- Quarterly Arm solution releases of integrated stack with CCA v1.1 features

Thank you!



Arm Solutions at Lightspeed