



Top-down Performance Analysis on Arm

openEuler Arm SIG Meetup

Yibo Cai
21-Jun-2024

What is Top-down Analysis?

- + To solve the pain of optimizing complex applications
- + Provides methodology to identify bottleneck on modern CPUs

A Top-Down Method for Performance Analysis and Counters Architecture

Ahmad Yasin

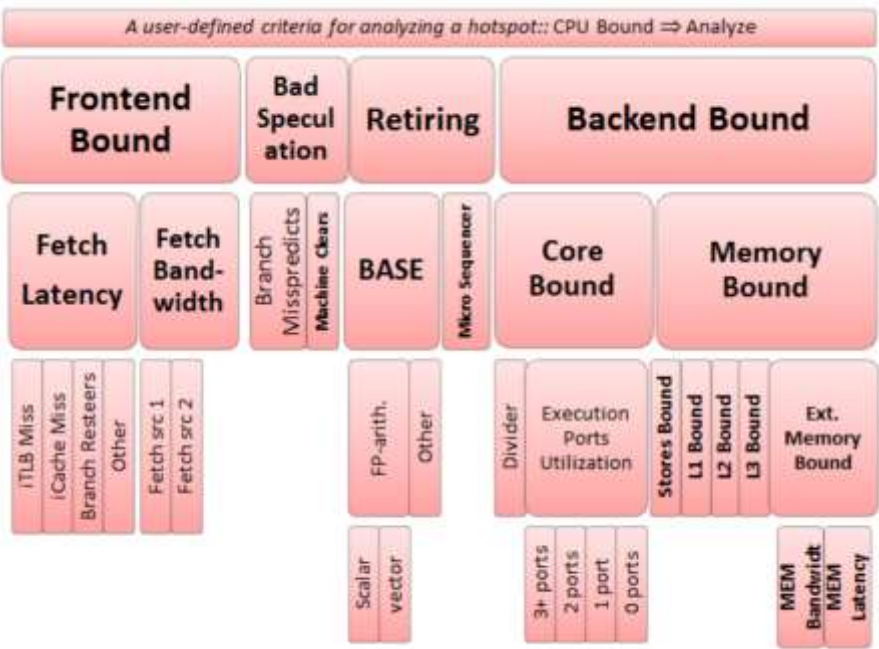
ahmad.yasin@intel.com

Intel Corporation, Architecture Group

Abstract

Optimizing an application's performance for a given microarchitecture has become painfully difficult. Increasing microarchitecture complexity, workload diversity, and the unmanageable volume of data produced by performance tools increase the optimization challenges. At the same time resource and time constraints get tougher with recently emerged segments. This further calls for accurate and prompt analysis methods.

In this paper a Top-Down Analysis is developed – a practical method to quickly identify true bottlenecks in out-of-order processors. The developed method uses designated performance counters in a structured hierarchical approach to quickly and, more importantly, correctly identify dominant



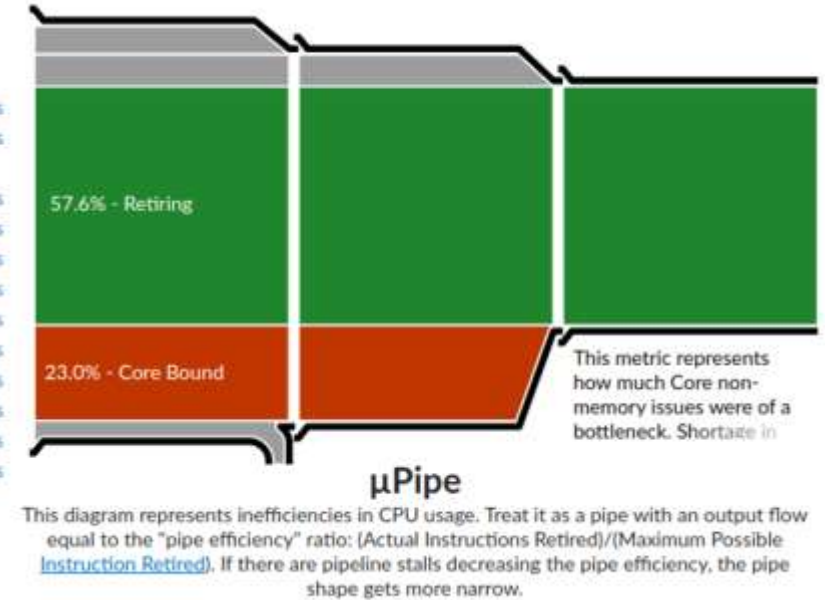
Top-down Analysis on x86

- + Multiple levels
- + Top level breakdown
 - *frontend_bound*
 - *backend_bound*
 - *bad_speculation*
 - *retiring*



- + Tooling: VTune, Toplev

| | |
|------------------------------|-------------------------|
| Elapsed Time: 10.429s | |
| Clockticks: | 28,209,500,000 |
| Instructions Retired: | 62,514,000,000 |
| CPI Rate: | 0.451 |
| Retiring: | 57.6% of Pipeline Slots |
| Light Operations: | 54.8% of Pipeline Slots |
| FP Arithmetic: | 0.0% of Pipeline Slots |
| Memory Operations: | 24.2% of Pipeline Slots |
| Fused Instructions: | 2.5% of Pipeline Slots |
| Non Fused Branches: | 1.6% of Pipeline Slots |
| Nop Instructions: | 0.4% of Pipeline Slots |
| Other: | 26.1% of Pipeline Slots |
| Heavy Operations: | 2.8% of Pipeline Slots |
| Front-End Bound: | 6.8% of Pipeline Slots |
| Bad Speculation: | 4.9% of Pipeline Slots |
| Back-End Bound: | 30.7% of Pipeline Slots |
| Memory Bound: | 7.7% of Pipeline Slots |
| L1 Bound: | 1.5% of Clockticks |
| L2 Bound: | 0.0% of Clockticks |
| L3 Bound: | 1.5% of Clockticks |
| DRAM Bound: | 4.4% of Clockticks |
| Store Bound: | 0.0% of Clockticks |
| Core Bound: | 23.0% of Pipeline Slots |
| Divider: | 4.4% of Clockticks |
| Port Utilization: | 15.7% of Clockticks |
| Cycles of 0 Ports Utilized: | 9.1% of Clockticks |
| Cycles of 1 Port Utilized: | 6.1% of Clockticks |
| Cycles of 2 Ports Utilized: | 9.1% of Clockticks |
| Cycles of 3+ Ports Utilized: | 30.3% of Clockticks |
| Vector Capacity Usage (FPU): | 12.5% |
| Average CPU Frequency: | 2.8 GHz |
| Total Thread Count: | 65 |
| Paused Time: | 0s |



Top-down Analysis on Arm – Level1

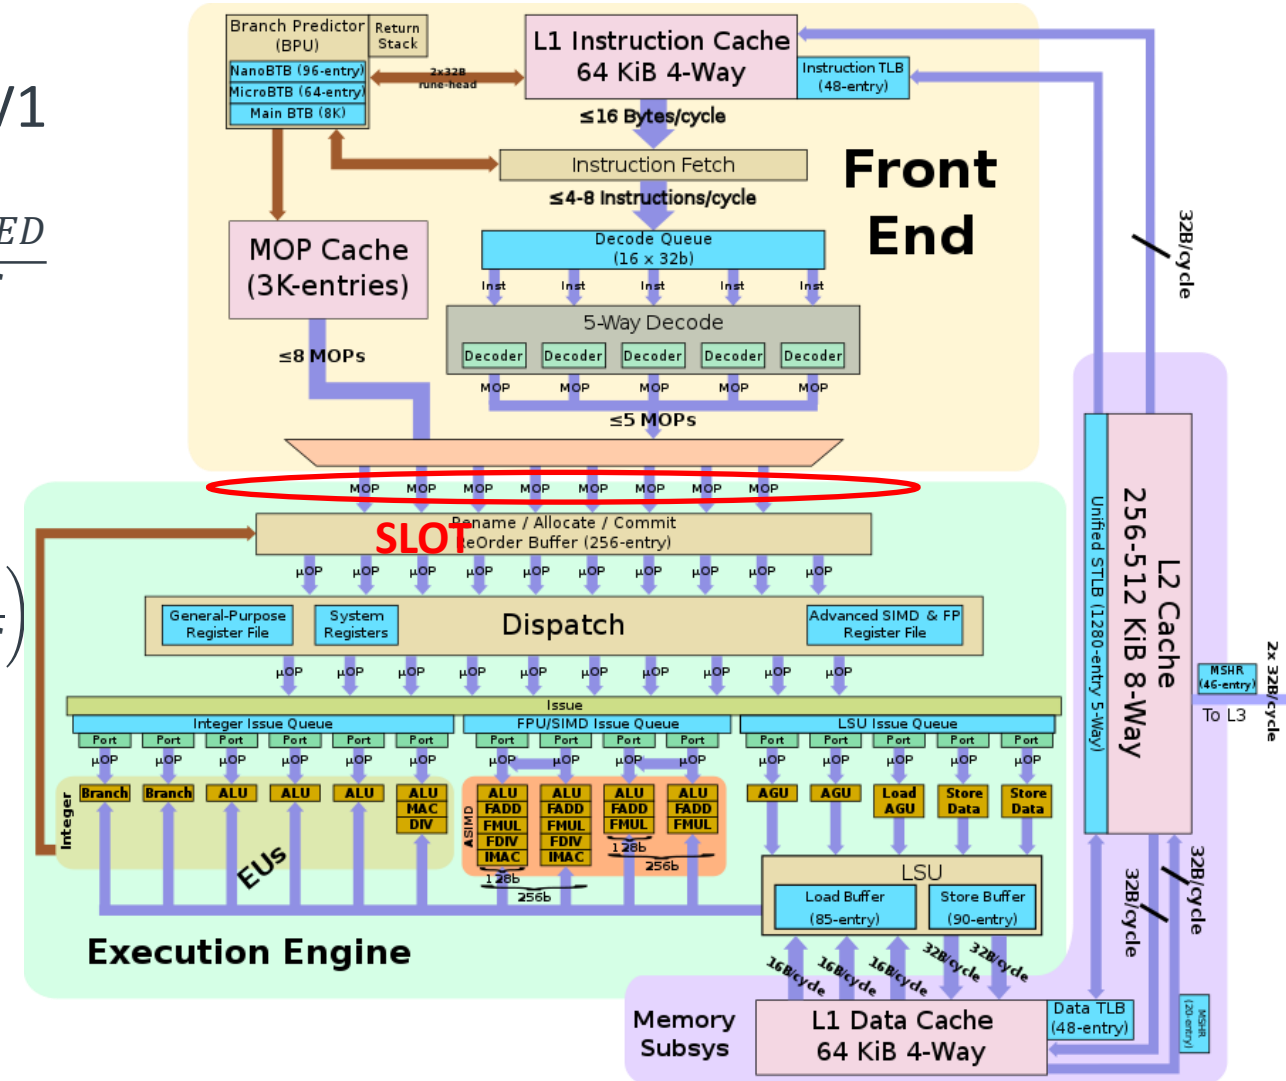
+ E.g., Top-down Level1 on Arm Neoverse-V1

$$frontend_bound = \frac{STALL_SLOT_FRONTEND}{8 * CPU_CYCLES} - \frac{4 * BR_MIS_PRED}{CPU_CYCLES}$$

$$backend_bound = \frac{STALL_SLOT_BACKEND}{8 * CPU_CYCLES}$$

$$bad_speculation = \left(1 - \frac{OP_RETIRED}{OP_SPEC}\right) * \left(1 - \frac{STALL_SLOT}{8 * CPU_CYCLES}\right) + \frac{4 * BR_MIS_PRED}{CPU_CYCLES}$$

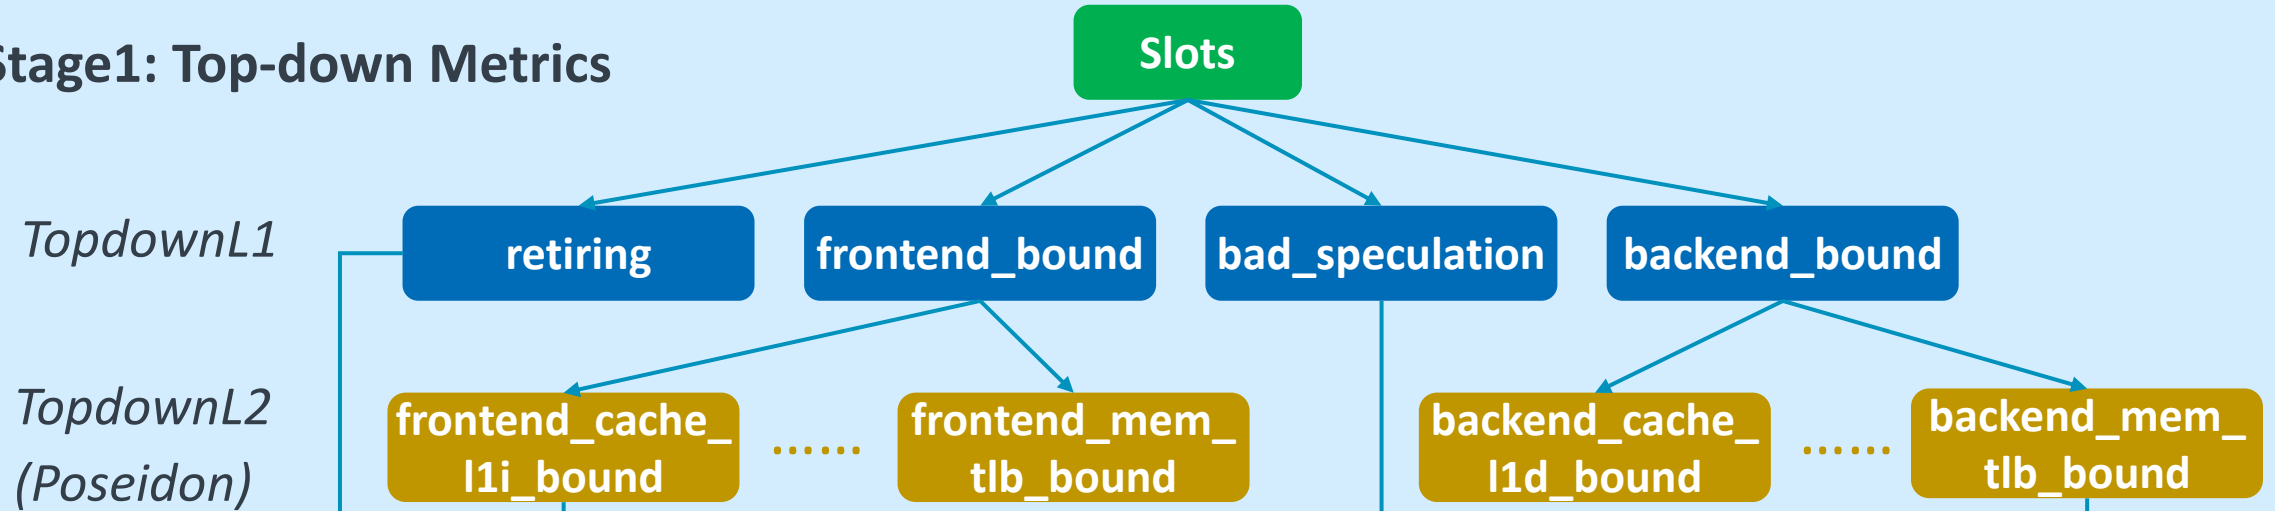
$$retiring = \frac{OP_RETIRED}{OP_SPEC} * \left(1 - \frac{STALL_SLOT}{8 * CPU_CYCLES}\right)$$



https://en.wikichip.org/wiki/arm_holdings/microarchitectures/neoverse_v1

Top-down Analysis on Arm – Two Stages

Stage1: Top-down Metrics



Stage2: μarch Metrics

OperationMix

- load_percentage
- store_percentage
- integer_dp_percentage
-

L1iCacheEffectiveness

- l1i_cache_mpki
- l1i_cache_miss_ratio

BranchEffectiveness

- branch_mpki
- branch_misprediction_ratio

DtlbEffectiveness

- dtlb_mpki
- dtlb_walk_ratio
-

Top-down Analysis on Arm – Tooling

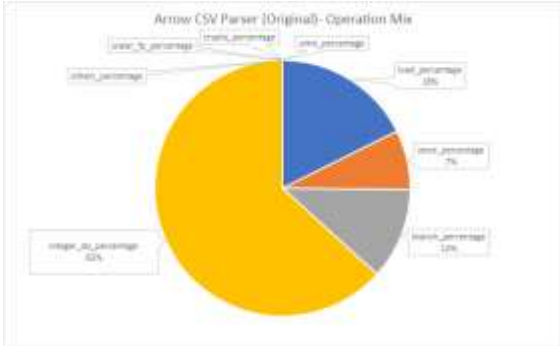
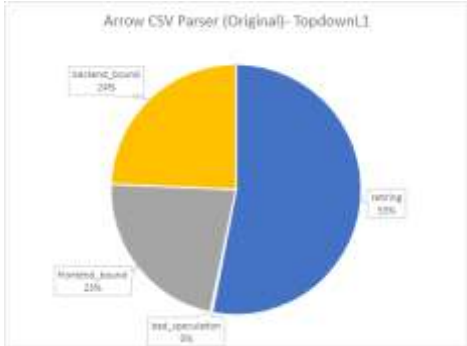
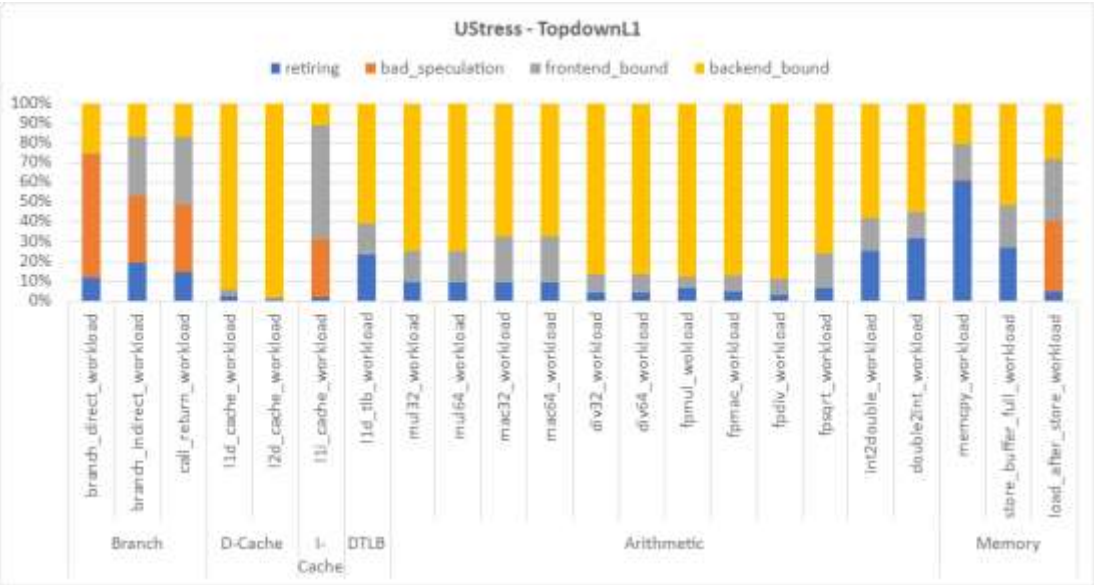
+ Linux perf tool

```
$ perf stat -M TopdownL1 -- workload
```

```
          31,678 BR_MIS_PRED          # 0.1 percent of slots bad_speculation
10,311,676,357 OP_SPEC                # 59.8 percent of slots retiring
10,336,289,342 STALL_SLOT
  3,441,336,410 CPU_CYCLES
10,297,795,725 OP_RETIRED
          30,386 BR_MIS_PRED          # 39.8 percent of slots frontend_bound
10,295,827,924 STALL_SLOT_FRONTEND
  3,441,227,523 CPU_CYCLES
          29,211 BR_MIS_PRED          # 0.2 percent of slots backend_bound
          40,563,280 STALL_SLOT_BACKEND
  3,438,935,083 CPU_CYCLES
```

Top-down Analysis on Arm – Tooling

+ Arm Telemetry Solution



Stage 1 (Top-down metrics)

[Topdown Level 1]

Frontend Bound..... 33.91% slots
Backend Bound..... 32.67% slots
Retiring..... 31.54% slots
Bad Speculation..... 0.95% slots

Stage 2 (uarch metrics)

[Branch Effectiveness]

(follows Frontend Bound)
(follows Bad Speculation)
Branch Misprediction Ratio..... 0.001 per branch
Branch MPKI..... 0.186 misses per 1,000 instructions

[Data TLB Effectiveness]

(follows Backend Bound)
DTLB MPKI..... 0.191 misses per 1,000 instructions
DTLB Walk Ratio..... 0.001 per TLB access

[Instruction TLB Effectiveness]

(follows Frontend Bound)
ITLB MPKI..... 0.008 misses per 1,000 instructions
ITLB Walk Ratio..... 0.000 per TLB access

[L1 Data Cache Effectiveness]

(follows Backend Bound)
L1D Cache Miss Ratio..... 0.020 per cache access
L1D Cache MPKI..... 6.930 misses per 1,000 instructions

..... omitted

[Operation Mix]

(follows Backend Bound)
(follows Retiring)
Speculative Branch Operations Percentage..... 14.52% instructions
Speculative Integer Operations Percentage..... 37.60% instructions
Speculative Load Operations Percentage..... 16.61% instructions
Speculative Floating Point Operations Percentage..... 0.00% instructions
Speculative SIMD Operations Percentage..... 23.07% instructions
Speculative Store Operations Percentage..... 9.33% instructions

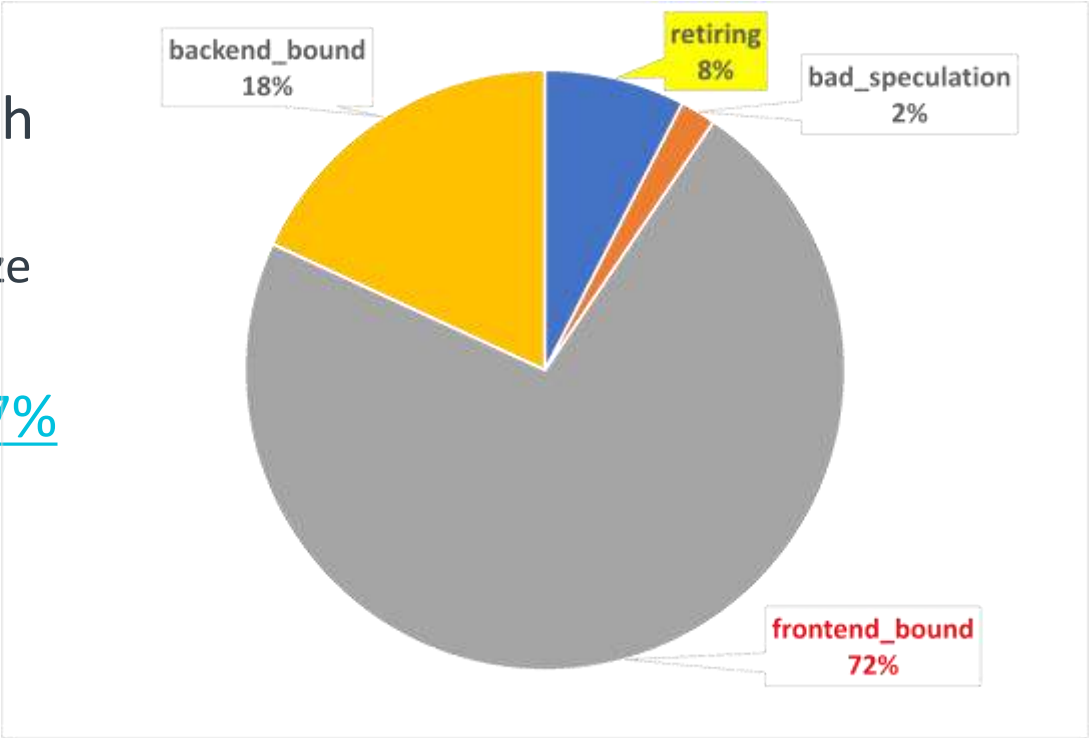
Top-down Analysis in Practice – MySQL

- + Topdown-tool shows MySQL write-only case suffers from significant frontend-bound and high instruction cache misses
 - PGO/AutoFDO/BOLT are common methods to optimize workloads with large code footprint

+ BOLT improves MySQL write performance by 17%

```
$ ./topdown-tool --perf-args "--tid <MYSQL_CLIENT_TID>"
Stage 1 (Topdown metrics)
=====
Frontend Bound..... 72.50% slots
Backend Bound..... 18.13% slots
Retiring..... 7.58% slots
Bad Speculation..... 1.93% slots

Stage 2 (uarch metrics)
=====
[L1 Instruction Cache Effectiveness]
L1I Cache Miss Ratio..... 0.194 per cache access
L1I Cache MPKI..... 64.789 misses per 1,000 instructions
```



| | Branch | L1d Cache | L1i Cache | L2 Cache |
|------|--------|-----------|-----------|----------|
| MPKI | 5.34 | 19.24 | 64.79 | 35.59 |
| Miss | 2.60% | 4.70% | 19.40% | 19.80% |

Top-down Analysis in Practice – Apache Arrow

- ✦ Arrow CSV Parser converts CSV data to Arrow columnar format

Name, City, Age

Mike, Shanghai, 15

Peter, Hangzhou, 14

Jack, Beijing, 17

➡

Name: "MikePeterJack", [0, 4, 9, 13]

City: "ShanghaiHangzhouBeijing", [0, 8, 16, 23]

Age: [15, 14, 17]

- ✦ Evaluate current performance

```
$ perf stat -e cycles,instructions \
    original/arrow-csv-parser-benchmark --benchmark_filter=ParseCSVVehiclesExample

-----
Benchmark                                     Time          CPU    Iterations UserCounters...
-----
ParseCSVVehiclesExample/iterations:200  10427125 ns   10424094 ns           200 bytes_per_second=1100.14M/s

Performance counter stats for 'original/arrow-csv-parser-benchmark --benchmark_filter=ParseCSVVehiclesExample'


    5454315340    cycles
    25288198650    instructions          #    4.64    insn per cycle

    2.106576442 seconds time elapsed

    2.086102000 seconds user
    0.020020000 seconds sys
```


Stage1 – Check the TopdownL1 Metrics

Arrow CSV Parser (Original) - TopdownL1



- + *Retiring* is above 50%, which matches the high IPC we observed.
- + No obvious bottleneck from *frontend_bound* and *backend_bound*.
- + For high *retiring* case, we may want to investigate the **Operation Mix** metrics.

Stage2 – Check the Operation Mix



- + Observed high volumes of integer calculation (*integer_dp_percentage* $\approx 2/3$)
- + Heavy integer calculation loading often means chances to leverage SIMD for better data parallelism

Why High *integer_dp_percentage*?

- ✦ The issue: special tokens (, \n, ...) hurts performance

| <i>Name, City, Age</i> | | |
|------------------------|---|-----------------------------------------------------|
| Mike, Shanghai, 15 | | <i>Name:</i> "MikePeterJack", [0,4,9,13] |
| Peter, Hangzhou, 14 | ➡ | <i>City:</i> "ShanghaiHangzhouBeijing", [0,8,16,23] |
| Jack, Beijing, 17 | | <i>Age:</i> [15, 14, 17] |

- ✦ Leverage Neon to process CSV data in batch

Mike, Shanghai, 15 Special tokens found. Go **slow** path. Check char by char.

Mike, Shanghai, 15 No special token. Go **fast** path. Process 8 chars at once.

```
bool Matches(uint8x8_t w) {  
    v = vceq_u8(w, vdup_n_u8('\r'));  
    v = vorr_u8(v, vceq_u8(w, vdup_n_u8('\n')));  
    v = vorr_u8(v, vceq_u8(w, delim_));  
    v = vorr_u8(v, vceq_u8(w, quote_));  
    v = vorr_u8(v, vceq_u8(w, escape_));  
    return (uint64_t)v != 0;  
}
```


Evaluate the Optimized Code

```
$ perf stat -e cycles,instructions \
    optimized/arrow-csv-parser-benchmark --benchmark_filter=ParseCSVVehiclesExample
```

```
-----
Benchmark                                Time            CPU    Iterations UserCounters...
-----
ParseCSVVehiclesExample/iterations:200    5642248 ns      5639743 ns      200 bytes_per_second=1.98576G/s
```

```
Performance counter stats for 'optimized/arrow-csv-parser-benchmark --benchmark_filter=ParseCSVVehiclesExample'
```



```
2971634240    cycles
12771507212    instructions          #    4.30  insn per cycle

1.154066440 seconds time elapsed

1.123869000 seconds user
0.027996000 seconds sys
```

- + Performance improves ~**80%** from 1.10G/s to 1.99G/s
- + Total instructions drops ~50% from 2.53E+10 to 1.28E+10

Compare the Operation Mix



- + *integer_dp_percentage* drops from $\sim 2/3$ to $\sim 1/3$
- + *simd_percentage* now occupies over $1/4$ total instructions



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks